

In [0]:

```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
```

In [0]:

```
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                           n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
```

In [3]:

```
X.shape, y.shape
```

Out[3]:

```
((50000, 15), (50000,))
```

In [0]:

```
from sklearn.model_selection import train_test_split
```

In [0]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)
```

In [6]:

```
X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

Out[6]:

```
((37500, 15), (37500,), (12500, 15), (12500,))
```

In [0]:

```
from sklearn import linear_model
```

In [8]:

```
# alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive' schedules.

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=15, penalty='l2',
                                tol=1e-3, verbose=2, learning_rate='constant')
clf
```

Out[8]:

```
SGDClassifier(alpha=0.0001, average=False, class_weight=None,
              early_stopping=False, epsilon=0.1, eta0=0.0001,
              fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
              loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
              penalty='l2', power_t=0.5, random_state=15, shuffle=True,
              tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

In [9]:

```
clf.fit(X=X_train, y=y_train)
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.02 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.03 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.04 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.05 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.05 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.06 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.07 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.08 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.08 seconds.
Convergence after 10 epochs took 0.08 seconds
```

Out[9]:

```
SGDClassifier(alpha=0.0001, average=False, class_weight=None,
              early_stopping=False, epsilon=0.1, eta0=0.0001,
              fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
              loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
              penalty='l2', power_t=0.5, random_state=15, shuffle=True,
              tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

In [10]:

```
clf.coef_, clf.coef_.shape, clf.intercept_
```

Out[10]:

```
(array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
          0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
          0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
 (1, 15),
 array([ -0.8531383]))
```

## Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

### Instructions

- Load the datasets(train and test) into the respective arrays
- Initialize the weight\_vector and intercept term randomly

- Calculate the initial log loss for the train and test data with the current weight and intercept and store it in a list
- for each epoch:
  - for each batch of data points in train: (keep batch size=1)
    - calculate the gradient of loss function w.r.t each weight in weight vector
    - Calculate the gradient of the intercept [check this](#)
    - Update weights and intercept (check the equation number 32 in the above mentioned [pdf](#)):
 
$$w^{(t+1)} \leftarrow (1 - \frac{\alpha}{N}) w^{(t)} + \alpha x_n (y_n - \sigma(w^{(t)T} x_n + b^{(t)}))$$

$$b^{(t+1)} \leftarrow (b^{(t)} + \alpha (y_n - \sigma(w^{(t)T} x_n + b^{(t)})))$$
    - calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
    - And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
    - append this loss in the list ( this will be used to see how loss is changing for each epoch after the training is over )
- Plot the train and test loss i.e on x-axis the epoch number, and on y-axis the loss
- **GOAL:** compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of  $10^{-3}$

In [0]:

```
import numpy as np
```

## Initializing W,B,eta0 and alpha

In [12]:

```
w = np.zeros_like(X_train[0])
b = 0
eta0 = 0.0001
alpha = 0.0001
N = len(X_train)
N
```

Out[12]:

37500

## Computing Log-loss

In [0]:

```
import math
# you can free to change all these codes/structure
def compute_log_loss(y_train, pred):
    sum1 = 0
    for i in range(len(pred)):
        sum1 += ((y_train[i] * math.log(pred[i])) + ((1-y_train[i]) * math.log((1-pred[i]))))

    loss = (-sum1/len(y_train))
    return loss
```

## Sigmoid and predict functions

In [14]:

```
def sigmoid(x,w,b):
    return (1/(1+np.exp(-(np.dot(x,w)+b))))

def predict(X_train,w,b):
    pred = []
```

```

for i in range(len(X_train)):
    pred.append(sigmoid(X_train[i],w,b))

return pred

train_loss = []
test_loss = []

#print(pred[:5])
train_pred = predict(X_train,w,b)
test_pred = predict(X_test,w,b)

#Computing log-loss
train_loss.append(compute_log_loss(y_train,train_pred))
test_loss.append(compute_log_loss(y_test,test_pred))
print(train_loss[0])
print(test_loss[0])

```

0.6931471805594285  
0.6931471805600672

## SGD with train and test loss

In [15]:

```

import random

for epoch in range(100):
    for i in range(N):
        batch = random.randrange(1,N)

        w = (( 1 - ( (alpha*eta0)/N) ) * w ) + ( (alpha*X_train[batch]) * ( y_train[batch] - sigmoid( X
_train[batch],w, b) ) )
        #w = (1 - ( (alpha * eta0)/N ) * w ) + ( ( alpha * X_train[batch] ) * ( y_train[batch] - sigmoi
d(X_train[batch],w,b ) ) )
        #b = b + (alpha * (y_train[batch] - sigmoid(X_train,w,b)))
        b = (b - ( alpha * ( -(y_train[batch]) + sigmoid(X_train[batch],w, b) ) ))

    y_train_ep = predict(X_train,w,b)
    y_test_ep = predict(X_test,w,b)

    train_loss.append(compute_log_loss(y_train,y_train_ep))
    test_loss.append(compute_log_loss(y_test,y_test_ep))

    print("Epoch",epoch,"train_loss",train_loss[-1:], 'test_loss',test_loss[-1:])

```

```

Epoch 0 train_loss [0.4037381695280155] test_loss [0.4051409460739654]
Epoch 1 train_loss [0.3881853608997544] test_loss [0.3898325942293472]
Epoch 2 train_loss [0.3829884680624984] test_loss [0.3850343141489612]
Epoch 3 train_loss [0.3810364322337181] test_loss [0.3829650524208779]
Epoch 4 train_loss [0.3793307474020105] test_loss [0.3811837756963011]
Epoch 5 train_loss [0.37899058749291276] test_loss [0.38093730513735213]
Epoch 6 train_loss [0.37860451654347943] test_loss [0.3809161630106351]
Epoch 7 train_loss [0.37859248967793074] test_loss [0.3806523284867069]
Epoch 8 train_loss [0.37839479771533124] test_loss [0.38062307300682996]
Epoch 9 train_loss [0.3785997381435557] test_loss [0.3801598326075643]
Epoch 10 train_loss [0.37875586516661724] test_loss [0.3808460189708544]
Epoch 11 train_loss [0.3785334950483176] test_loss [0.38043706353596246]
Epoch 12 train_loss [0.378319311341773] test_loss [0.38046905269524517]
Epoch 13 train_loss [0.37849586907479216] test_loss [0.38106104382028905]
Epoch 14 train_loss [0.37844457709883084] test_loss [0.38077869343296433]
Epoch 15 train_loss [0.37828723193219166] test_loss [0.38057143966617346]
Epoch 16 train_loss [0.3782556257651049] test_loss [0.3802836106482321]
Epoch 17 train_loss [0.37828596979548795] test_loss [0.38025606448603394]
Epoch 18 train_loss [0.378790052102815] test_loss [0.3806353954847428]
Epoch 19 train_loss [0.3785487186734486] test_loss [0.3806716857464192]

```

Epoch 20	train_loss	[0.37834039873327374]	test_loss	[0.380488759085167]
Epoch 21	train_loss	[0.37856679503859286]	test_loss	[0.3808323157183337]
Epoch 22	train_loss	[0.3784557350099683]	test_loss	[0.3805610330744679]
Epoch 23	train_loss	[0.3782690945458786]	test_loss	[0.3801988431722665]
Epoch 24	train_loss	[0.37833664657542676]	test_loss	[0.3805911635537753]
Epoch 25	train_loss	[0.3784389083729466]	test_loss	[0.3804360786727016]
Epoch 26	train_loss	[0.3784692967084271]	test_loss	[0.38015570499625095]
Epoch 27	train_loss	[0.37841150141590435]	test_loss	[0.3806889376559375]
Epoch 28	train_loss	[0.37831904368221814]	test_loss	[0.38042784510305055]
Epoch 29	train_loss	[0.37826504783083614]	test_loss	[0.3801105079157159]
Epoch 30	train_loss	[0.3784010305318321]	test_loss	[0.3802748021726142]
Epoch 31	train_loss	[0.37868579423054605]	test_loss	[0.3806650095170454]
Epoch 32	train_loss	[0.37839632632599973]	test_loss	[0.38008283843529606]
Epoch 33	train_loss	[0.37860429504902243]	test_loss	[0.38096074308666156]
Epoch 34	train_loss	[0.3783298607628876]	test_loss	[0.3804859460392227]
Epoch 35	train_loss	[0.37824503502500256]	test_loss	[0.38015237314926836]
Epoch 36	train_loss	[0.37879789758874044]	test_loss	[0.3803794993881466]
Epoch 37	train_loss	[0.3785472200787773]	test_loss	[0.38040218673309906]
Epoch 38	train_loss	[0.37838983389721187]	test_loss	[0.3802951235580614]
Epoch 39	train_loss	[0.378407534818248]	test_loss	[0.38084593052827237]
Epoch 40	train_loss	[0.37849607363090726]	test_loss	[0.38018167722818375]
Epoch 41	train_loss	[0.37887814382749574]	test_loss	[0.3805813042564888]
Epoch 42	train_loss	[0.37844404692467415]	test_loss	[0.38050030006947005]
Epoch 43	train_loss	[0.3786995542218357]	test_loss	[0.38059106130396336]
Epoch 44	train_loss	[0.3782801139447077]	test_loss	[0.3806497539785946]
Epoch 45	train_loss	[0.3785891608233322]	test_loss	[0.3808594095194392]
Epoch 46	train_loss	[0.37838072581047594]	test_loss	[0.3804360100608189]
Epoch 47	train_loss	[0.3781714066728731]	test_loss	[0.38016317734765137]
Epoch 48	train_loss	[0.37828268136228665]	test_loss	[0.38043582085017064]
Epoch 49	train_loss	[0.3783174939444274]	test_loss	[0.38047780731136643]
Epoch 50	train_loss	[0.37842365245193715]	test_loss	[0.3803371019224976]
Epoch 51	train_loss	[0.3783860620355435]	test_loss	[0.3801683995299245]
Epoch 52	train_loss	[0.37838055985588154]	test_loss	[0.3803110665815696]
Epoch 53	train_loss	[0.3783420480114288]	test_loss	[0.380476787936208]
Epoch 54	train_loss	[0.37865727386751236]	test_loss	[0.3804228973577421]
Epoch 55	train_loss	[0.3782474839632092]	test_loss	[0.3804848064206839]
Epoch 56	train_loss	[0.3782923586710222]	test_loss	[0.3804011586827127]
Epoch 57	train_loss	[0.378460508474135]	test_loss	[0.380454314739438]
Epoch 58	train_loss	[0.3785546435055251]	test_loss	[0.38099400105225767]
Epoch 59	train_loss	[0.3783692168838415]	test_loss	[0.3808645734690277]
Epoch 60	train_loss	[0.3790874604560089]	test_loss	[0.38078547195640866]
Epoch 61	train_loss	[0.378385870247855]	test_loss	[0.38066898640585073]
Epoch 62	train_loss	[0.37830267246996796]	test_loss	[0.38053803311555623]
Epoch 63	train_loss	[0.37824347026778515]	test_loss	[0.38021824983708474]
Epoch 64	train_loss	[0.3784408428316232]	test_loss	[0.38043835737265297]
Epoch 65	train_loss	[0.37836581233019456]	test_loss	[0.38068339216397623]
Epoch 66	train_loss	[0.37845100227737377]	test_loss	[0.3807824798604219]
Epoch 67	train_loss	[0.37847992231806377]	test_loss	[0.3804682657252031]
Epoch 68	train_loss	[0.3784411767775674]	test_loss	[0.38092653904334467]
Epoch 69	train_loss	[0.3782921860759299]	test_loss	[0.38020459935435513]
Epoch 70	train_loss	[0.3785881508556596]	test_loss	[0.3809583408438084]
Epoch 71	train_loss	[0.3783378803963715]	test_loss	[0.3803961820293535]
Epoch 72	train_loss	[0.37826974638864536]	test_loss	[0.38034150901414987]
Epoch 73	train_loss	[0.378530156100854]	test_loss	[0.3807972266983999]
Epoch 74	train_loss	[0.37858045151545444]	test_loss	[0.3805177474743326866]
Epoch 75	train_loss	[0.37820712538405743]	test_loss	[0.38036716256505965]
Epoch 76	train_loss	[0.3784912937264594]	test_loss	[0.38031813672802245]
Epoch 77	train_loss	[0.3784848584366938]	test_loss	[0.38055749435134995]
Epoch 78	train_loss	[0.3783430587083698]	test_loss	[0.3807588672280358]
Epoch 79	train_loss	[0.37837175280467616]	test_loss	[0.38047390286230287]
Epoch 80	train_loss	[0.3784321673710929]	test_loss	[0.3807268129664003]
Epoch 81	train_loss	[0.3788426523451305]	test_loss	[0.3811161603145213]
Epoch 82	train_loss	[0.3787392300059855]	test_loss	[0.3810545793385992]
Epoch 83	train_loss	[0.37859481559275476]	test_loss	[0.38074771166607935]
Epoch 84	train_loss	[0.37836201947462694]	test_loss	[0.3805159208115878]
Epoch 85	train_loss	[0.3784248806056946]	test_loss	[0.3805133354351515]
Epoch 86	train_loss	[0.3786555140887769]	test_loss	[0.3814232068641636]
Epoch 87	train_loss	[0.37863614401627127]	test_loss	[0.3812271984515548]
Epoch 88	train_loss	[0.3785168930344105]	test_loss	[0.38013927520184143]
Epoch 89	train_loss	[0.3786452168124955]	test_loss	[0.3808409840228766]
Epoch 90	train_loss	[0.3786334208030699]	test_loss	[0.38057041732211777]
Epoch 91	train_loss	[0.3785733640230413]	test_loss	[0.38056684276704683]
Epoch 92	train_loss	[0.3784738321160625]	test_loss	[0.3809593605655751]
Epoch 93	train_loss	[0.37839491220362464]	test_loss	[0.3807692128526246]
Epoch 94	train_loss	[0.37847350866712937]	test_loss	[0.38017437066722504]
Epoch 95	train_loss	[0.3783669260998517]	test_loss	[0.3805901461418767]
Epoch 96	train_loss	[0.37831901692810516]	test_loss	[0.3803873096475521]

```
Epoch 96 train_loss [0.37850949084546265] test_loss [0.38023338233336923]
Epoch 97 train_loss [0.37850949084546265] test_loss [0.38023338233336923]
Epoch 98 train_loss [0.3782134534745643] test_loss [0.3802217437863191]
Epoch 99 train_loss [0.3784287208870628] test_loss [0.3803246990204749]
```

## Plot of train-loss and test-loss

In [16]:


```
import matplotlib.pyplot as plt

#plt.figure(figsize=(9,6))

plt.plot(train_loss[1:], label="train log-loss")
plt.plot(test_loss[1:], label="test log-loss")

plt.xlabel("Epoch")
plt.ylabel("Log-loss")
plt.legend()
plt.title("Log-loss")

plt.show()
```



## Weight and intercept

In [17]:

```
print("Weight vector(W) :",w)
print('Intercept(B) : ',b)
```

```
Weight vector(W) : [-0.43026846  0.19244311 -0.14315103  0.34086735 -0.22428848  0.56679028
 -0.44654246 -0.09243658  0.22266273  0.17870868  0.20592653 -0.00220059
 -0.0813303   0.33423687  0.03268197]
Intercept(B) : -0.891685729814871
```

## Diff btw skcit learn and custom implemented weights

In [18]:

```
# these are the results we got after we implemented sgd and found the optimal weights and intercept
w-clf.coef_, b-clf.intercept_
```

Out[18]:

```
(array([[ -0.00690154,  0.00696746,  0.00543933, -0.00057672, -0.01610178,
          0.0066245 ,  0.00588236,  0.00165155,  0.01338953, -0.00213258,
          0.00887463, -0.00641975, -0.0017266 , -0.00429115,  0.01001476]]),
 array([ -0.03854743]))
```

## Calculating accuracy

In [19]:

```
def pred(w,b, X):
    N = len(X)
    predict = []
    for i in range(N):
        if 1/(1+np.exp(-(np.dot(X[i],w)+b))) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
            predict.append(1)
        else:
            predict.append(0)
    return np.array(predict)

print("Train accuracy : ",1-np.sum(y_train - pred(w,b,X_train))/len(X_train))
print("Test accuracy :",1-np.sum(y_test - pred(w,b,X_test))/len(X_test))
```

Train accuracy : 0.9541866666666666

Test accuracy : 0.95256